

Data Analytics with Cognos-

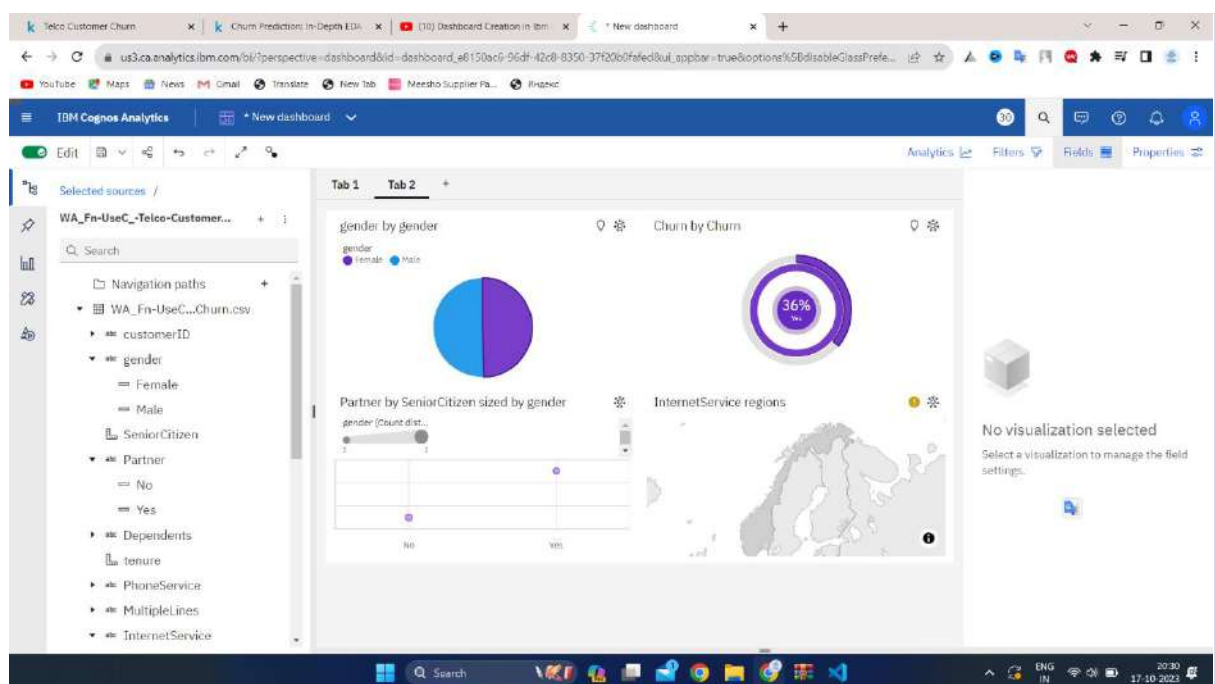
Group2

Project Title: Customer Churn Prediction

Group Members;

- 1) Karthika G (TL)
- 2) Poornima
- 3) Sindhu V
- 4) Varalakshmi M
- 5) Aravind M

IBM Cognos Visualization:-



1.0 Business Understanding

CUSTOMER CHURN PREDICTION

1.1 Background

Acquiring customers in the telecommunications industry is a challenging task due to intense competition, evolving technology, high acquisition costs and customer retention concerns. Telecommunication companies need to stand out by offering specialized services, stay updated with technology and making sure they don't spend too much on customer acquisition. Among the three main revenue-generating strategies (acquiring new customers, upselling to existing customers, and increasing customer retention), increasing retention has proven to be less costly and the most profitable. To achieve this, reducing customer churn, the movement from one provider to another, becomes a critical focus in the highly competitive telecommunications sector. By prioritizing customer retention and addressing churn, telecom companies can maximize profitability and long-term success. Therefore, finding those factors that increase customer churn is important to take necessary actions to reduce this churn. The main goal of our project is to develop an understanding of the cause of customer churn which assists telecom operators to predict customers who are most likely subject to churn, and what to do to retain the most valuable customer.

1.2 Problem Statement

SyriaTel Telecommunications has experienced a substantial increase in customer churn rates in American states during the last financial period, resulting in a significant number of customers switching to competitors. Recognizing the urgency of understanding the underlying factors contributing to this trend, the marketing department at SyriaTel has taken proactive measures. They have engaged a consortium of scientists to develop a predictive model capable of identifying customers who are likely to churn, as well as analyzing their behaviors. This initiative aims to address the pressing challenge of customer attrition, which poses a threat to SyriaTel's bottom line and overall revenue growth. By leveraging the insights provided by the predictive model, SyriaTel intends to implement targeted strategies that will enhance customer retention, safeguard their bottom line, and foster new avenues for revenue growth.

1.3 Objectives

- To understand which factors or variables contribute the most to customer churn.
- To identify different customer segments based on churn behaviour
- To develop a model that can accurately predict customer churn.
- To obtain valuable insights that help generate the best recommendations to protect Syriatel's revenue.

2.0 Data Understanding

The dataset contains various features related to telecom customer behavior, service usage, and account information. It includes details such as the customer's state, account length, area code, phone number, international plan, voice mail plan, number of voicemail messages, and the total duration and charges for calls made during the day, evening, and night. It also includes information on international calls, customer service calls, and whether or not the customer churned (terminated their contract). The data will be suitable to build a classifier to predict whether a customer will ("soon") stop doing business with SyriaTel telecommunications company.

Data Source : <https://www.kaggle.com/datasets/becksddf/churn-in-telecoms-dataset>
(<https://www.kaggle.com/datasets/becksddf/churn-in-telecoms-dataset>)

Summary of Features in the Dataset

- **state:** the state the customer lives in - seems like American states
- **account length:** the number of days the customer has had an account
- **area code:** the area code of the customer
- **phone number:** the phone number of the customer
- **international plan:** true if the customer has the international plan, otherwise false
- **voice mail plan:** true if the customer has the voice mail plan, otherwise false
- **number vmail messages:** the number of voicemails the customer has sent
- **total day minutes:** total number of minutes the customer has been in calls during the day
- **total day calls:** total number of calls the user has done during the day
- **total day charge:** total amount of money the customer was charged by the Telecom company for calls during the day
- **total eve minutes:** total number of minutes the customer has been in calls during the evening
- **total eve calls:** total number of calls the customer has done during the evening
- **total eve charge:** total amount of money the customer was charged by the Telecom company for calls during the evening
- **total night minutes:** total number of minutes the customer has been in calls during the night
- **total night calls:** total number of calls the customer has done during the night
- **total night charge:** total amount of money the customer was charged by the Telecom company for calls during the night
- **total intl minutes:** total number of minutes the user has been in international calls
- **total intl calls:** total number of international calls the customer has done
- **total intl charge:** total amount of money the customer was charged by the Telecom company for international calls
- **customer service calls:** number of calls the customer has made to customer service
- **churn:** target variable which is true if the customer terminated their contract, otherwise false

```

#import all the libraries

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV,
cross_val_score, RandomizedSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report,
explained_variance_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, precision_recall_curve,
recall_score, precision_score, f1_score, accuracy_score, roc_curve,
roc_auc_score, auc
from xgboost import XGBClassifier, plot_importance
from sklearn.utils import resample
from imblearn.over_sampling import SMOTE
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)

```

```

#load the dataset and view the first 5 columns

```

```

churn_df = pd.read_csv("syriatel_data.csv")
churn_df.head(5)

```

	state	account length	area	code	phone number	international	plan \
0	KS	128		415	382-4657		no
1	OH	107		415	371-7191		no
2	NJ	137		415	358-1921		no
3	OH	84		408	375-9999		yes
4	OK	75		415	330-6626		yes

	voice mail plan	number vmail	messages	total day	minutes	total day
calls \						
0	yes		25		265.1	
110						
1	yes		26		161.6	
123						
2	no		0		243.4	
114						
3	no		0		299.4	
71						
4	no		0		166.7	
113						

	total day charge	...	total eve calls	total eve charge	\
0	45.07	...	99	16.78	
1	27.47	...	103	16.62	
2	41.38	...	110	10.30	
3	50.90	...	88	5.26	
4	28.34	...	122	12.61	

	total night minutes	total night calls	total night charge	\
0	244.7	91	11.01	
1	254.4	103	11.45	
2	162.6	104	7.32	
3	196.9	89	8.86	
4	186.9	121	8.41	

	total intl minutes	total intl calls	total intl charge	\
0	10.0	3	2.70	
1	13.7	3	3.70	
2	12.2	5	3.29	
3	6.6	7	1.78	
4	10.1	3	2.73	

	customer service calls	churn
0	1	False
1	1	False
2	0	False
3	2	False
4	3	False

[5 rows x 21 columns]

view the last 5 columns to check for any differences
 churn_df.tail(5)

	state	account length	area code	phone number	international plan
3328	AZ	192	415	414-4276	no
3329	WV	68	415	370-3271	no
3330	RI	28	510	328-8230	no
3331	CT	184	510	364-6381	yes
3332	TN	74	415	400-4344	no

	voice mail plan	number vmail messages	total day minutes	\
3328	yes	36	156.2	
3329	no	0	231.1	
3330	no	0	180.8	
3331	no	0	213.8	

3332	yes	25	234.4
------	-----	----	-------

	total day calls	total day charge	...	total eve calls	\
3328	77	26.55	...	126	
3329	57	39.29	...	55	
3330	109	30.74	...	58	
3331	105	36.35	...	84	
3332	113	39.85	...	82	

	total eve charge	total night minutes	total night calls	\
3328	18.32	279.1	83	
3329	13.04	191.3	123	
3330	24.55	191.9	91	
3331	13.57	139.2	137	
3332	22.60	241.4	77	

	total night charge	total intl minutes	total intl calls	\
3328	12.56	9.9	6	
3329	8.61	9.6	4	
3330	8.64	14.1	6	
3331	6.26	5.0	10	
3332	10.86	13.7	4	

	total intl charge	customer service calls	churn
3328	2.67	2	False
3329	2.59	3	False
3330	3.81	2	False
3331	1.35	2	False
3332	3.70	0	False

[5 rows x 21 columns]

#checking for column features

churn_df.columns

```
Index(['state', 'account length', 'area code', 'phone number',
      'international plan', 'voice mail plan', 'number vmail
messages',
      'total day minutes', 'total day calls', 'total day charge',
      'total eve minutes', 'total eve calls', 'total eve charge',
      'total night minutes', 'total night calls', 'total night
charge',
      'total intl minutes', 'total intl calls', 'total intl charge',
      'customer service calls', 'churn'],
      dtype='object')
```

#concise summary of the dataset

churn_df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account length                      3333 non-null   int64
2   area code                          3333 non-null   int64
3   phone number                       3333 non-null   object
4   international plan                 3333 non-null   object
5   voice mail plan                    3333 non-null   object
6   number vmail messages              3333 non-null   int64
7   total day minutes                  3333 non-null   float64
8   total day calls                    3333 non-null   int64
9   total day charge                   3333 non-null   float64
10  total eve minutes                  3333 non-null   float64
11  total eve calls                    3333 non-null   int64
12  total eve charge                   3333 non-null   float64
13  total night minutes                3333 non-null   float64
14  total night calls                  3333 non-null   int64
15  total night charge                 3333 non-null   float64
16  total intl minutes                 3333 non-null   float64
17  total intl calls                   3333 non-null   int64
18  total intl charge                  3333 non-null   float64
19  customer service calls             3333 non-null   int64
20  churn                              3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB

```

#shape of the dataset

```
churn_df.shape
```

```
(3333, 21)
```

summary statistics for numerical columns

```
churn_df.describe()
```

	account length	area code	number vmail messages	total day minutes
count	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	437.182418	8.099010	179.775098
std	39.822106	42.371290	13.688365	54.467389
min	1.000000	408.000000	0.000000	0.000000
25%	74.000000	408.000000	0.000000	

143.700000

50%	101.000000	415.000000	0.000000
-----	------------	------------	----------

179.400000

75%	127.000000	510.000000	20.000000
-----	------------	------------	-----------

216.400000

max	243.000000	510.000000	51.000000
-----	------------	------------	-----------

350.800000

	total day calls	total day charge	total eve minutes	total eve
--	-----------------	------------------	-------------------	-----------

calls \

count	3333.000000	3333.000000	3333.000000
-------	-------------	-------------	-------------

3333.000000

mean	100.435644	30.562307	200.980348
------	------------	-----------	------------

100.114311

std	20.069084	9.259435	50.713844
-----	-----------	----------	-----------

19.922625

min	0.000000	0.000000	0.000000
-----	----------	----------	----------

0.000000

25%	87.000000	24.430000	166.600000
-----	-----------	-----------	------------

87.000000

50%	101.000000	30.500000	201.400000
-----	------------	-----------	------------

100.000000

75%	114.000000	36.790000	235.300000
-----	------------	-----------	------------

114.000000

max	165.000000	59.640000	363.700000
-----	------------	-----------	------------

170.000000

	total eve charge	total night minutes	total night calls \
--	------------------	---------------------	---------------------

count	3333.000000	3333.000000	3333.000000
-------	-------------	-------------	-------------

mean	17.083540	200.872037	100.107711
------	-----------	------------	------------

std	4.310668	50.573847	19.568609
-----	----------	-----------	-----------

min	0.000000	23.200000	33.000000
-----	----------	-----------	-----------

25%	14.160000	167.000000	87.000000
-----	-----------	------------	-----------

50%	17.120000	201.200000	100.000000
-----	-----------	------------	------------

75%	20.000000	235.300000	113.000000
-----	-----------	------------	------------

max	30.910000	395.000000	175.000000
-----	-----------	------------	------------

	total night charge	total intl minutes	total intl calls \
--	--------------------	--------------------	--------------------

count	3333.000000	3333.000000	3333.000000
-------	-------------	-------------	-------------

mean	9.039325	10.237294	4.479448
------	----------	-----------	----------

std	2.275873	2.791840	2.461214
-----	----------	----------	----------

min	1.040000	0.000000	0.000000
-----	----------	----------	----------

25%	7.520000	8.500000	3.000000
-----	----------	----------	----------

50%	9.050000	10.300000	4.000000
-----	----------	-----------	----------

75%	10.590000	12.100000	6.000000
-----	-----------	-----------	----------

max	17.770000	20.000000	20.000000
-----	-----------	-----------	-----------

	total intl charge	customer service calls
--	-------------------	------------------------

count	3333.000000	3333.000000
-------	-------------	-------------

mean	2.764581	1.562856
------	----------	----------

std	0.753773	1.315491
min	0.000000	0.000000
25%	2.300000	1.000000
50%	2.780000	1.000000
75%	3.270000	2.000000
max	5.400000	9.000000

Observations

- Our data consists of 3333 rows and 21 columns
- Data has both continuous and categorical features comprising of the following data types; objects, integers, float and booleans
- Churn which is our target variable is of data type boolean.
- We can also see statistical summary of the numerical records based on their count, median, mean, standard deviation, percentiles, minimum and maximum values.

3.0 Data Preparation

3.1 Data Cleaning

This section prepares the data for EDA and modeling. The dataset will be checked for:

- duplicated rows
- missing values
- In our analysis, we will drop phone numbers as they do not provide any relevant insights.
- Create two variables for our numerical and categorical data types respectively

```
# check for duplicate records
```

```
churn_df.duplicated().sum()
```

```
0
```

```
# check missing values
```

```
churn_df.isnull().sum()
```

```
state          0
account length 0
area code      0
phone number   0
international plan 0
voice mail plan 0
number vmail messages 0
total day minutes 0
total day calls 0
total day charge 0
total eve minutes 0
```

```

total eve calls      0
total eve charge     0
total night minutes  0
total night calls    0
total night charge   0
total intl minutes   0
total intl calls     0
total intl charge    0
customer service calls 0
churn                0
dtype: int64

```

Dropping phone number because it lacks information about customer behavior.

```
churn_df = churn_df.drop('phone number', axis=1)
```

confirming we have dropped phone number

```
churn_df.columns
```

```

Index(['state', 'account length', 'area code', 'international plan',
       'voice mail plan', 'number vmail messages', 'total day
minutes',
       'total day calls', 'total day charge', 'total eve minutes',
       'total eve calls', 'total eve charge', 'total night minutes',
       'total night calls', 'total night charge', 'total intl
minutes',
       'total intl calls', 'total intl charge', 'customer service
calls',
       'churn'],
      dtype='object')

```

Categorical and numerical variables

```
cat_vars = []
```

```
num_vars = []
```

```

for col in churn_df.columns:
    if churn_df[col].dtype == 'object':
        cat_vars.append(col)
    else:
        num_vars.append(col)
num_vars.pop(-1)
print("-----")
print('Categorical variables:', cat_vars)
print("-----")
print('Numerical variables:', num_vars)
print("-----")

```

```
-----  
Categorical variables: ['state', 'international plan', 'voice mail  
plan']  
-----
```

```
Numerical variables: ['account length', 'area code', 'number vmail  
messages', 'total day minutes', 'total day calls', 'total day charge',  
'total eve minutes', 'total eve calls', 'total eve charge', 'total  
night minutes', 'total night calls', 'total night charge', 'total intl  
minutes', 'total intl calls', 'total intl charge', 'customer service  
calls']  
-----
```

3.2 Univariate Analysis

a. Analysis of the target variable "Churn"

```
#Value_count of target variable
```

```
churn_df["churn"].value_counts()
```


```
False    2850
```

```
True      483
```

```
Name: churn, dtype: int64
```

```
#count visualization
```

```
churn_df['churn'].value_counts().plot(  
    kind='bar',color=['blue', 'orange']).set_title(  
    "Distribution of customers who Churned and those who didn't");
```



#Check percent of current customers that have churned (True) and those that didn't (False)

```
churn_df["churn"].value_counts(normalize=True) * 100
```

```
False    85.508551
```

```
True     14.491449
```

```
Name: churn, dtype: float64
```


#To get the pie chart to analyze churn

```
churn_df ['churn'].value_counts().plot.pie(explode=[0.05,0.05],
```

```
autopct='%1.1f%%', startangle=90,shadow=True, figsize=(8,8))
```

```
plt.title('Pie Chart for Churn')
```

```
plt.show()
```



The above visualization shows customers with active contracts and those that have terminated.

When we check our target variable "churn" it indicates that the majority, which is 85% of the customers in the churn_df dataset are active and the rest, which is about 14.5% are inactive.

This means that the dataset primarily consists of active customers, with a relatively smaller portion of inactive customers presenting a case of class imbalance which will require appropriate strategies to handle before modeling as an imbalanced class can cause the model to make false predictions.

















b. Univariate analysis for Numerical Variables

```
# Create subplots for each numerical variable  
num_plots = len(num_vars)
```

```
num_rows = 4
num_cols = 4
fig, axes = plt.subplots(nrows=num_rows, ncols=num_cols, figsize=(20,
30))

for i, var in enumerate(num_vars):
    row = i // num_cols
    col = i % num_cols
    axes[row, col].hist(churn_df[var], bins=30)
    axes[row, col].set_title(var, fontsize=20)
    axes[row, col].set_xlabel('Value', fontsize=20)
    axes[row, col].set_ylabel('Frequency', fontsize=20)

plt.tight_layout()
plt.show()
```



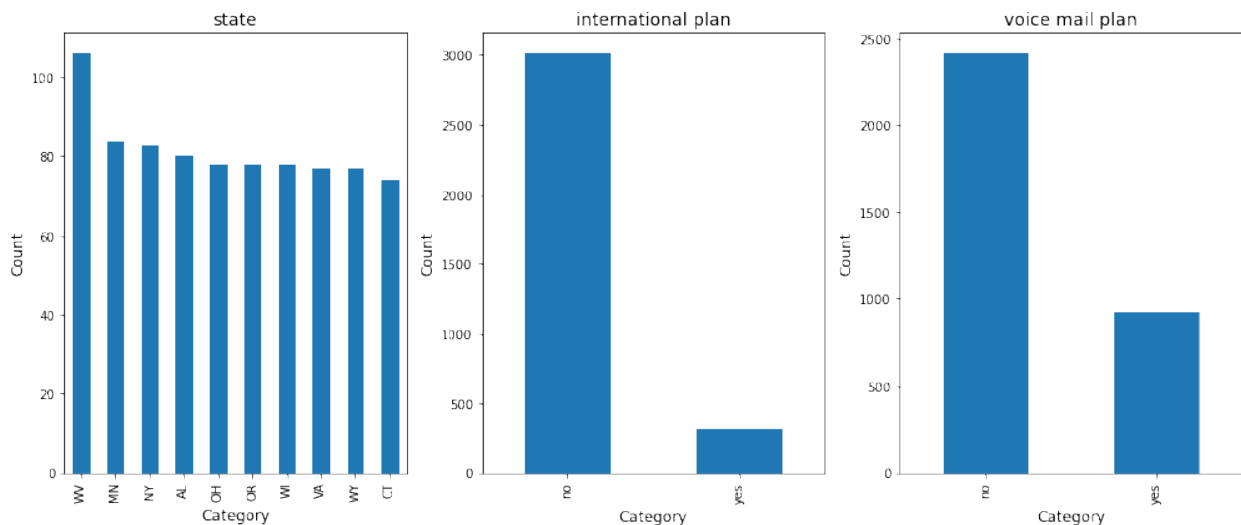
- The majority of the features in the data exhibit a normal distribution. This characteristic implies that the data points within these features tend to cluster around the mean, with relatively fewer occurrences of extreme values.
- Majority of customers in the dataset have made one customer service call.
- The highest number of calls made to customer service is 9 calls.
- The total international calls and customer service calls are skewed to the right

c. Univariate analysis for Categorical Variables

```
fig, ax = plt.subplots(nrows=1, ncols=3, figsize=(14, 6))

for i, cat_var in enumerate(cat_vars):
    top_ten_cats = churn_df[cat_var].value_counts().nlargest(10)
    top_ten_cats.plot(kind='bar', ax=ax[i])
    ax[i].set_title(cat_var, fontsize=14)
    ax[i].set_xlabel('Category', fontsize=12)
    ax[i].set_ylabel('Count', fontsize=12)

plt.tight_layout()
plt.show()
```



- The top five American states that syriatel operates in are West Virginia, Minnesota, New York, Alabama and Oregon respectively.
- The majority of customers in the dataset do not have an international plan or a voice mail plan

3.3 Bivariate Analysis

a. Analysis of churned Customers based on International Plan

#Churned customers by international plan

```
churn_international_plan = churn_df.groupby("international plan")
```




```
[ "churn"].sum().reset_index()
churn_international_plan
```

```
international plan churn
0                no   346
1                yes  137
```

Lets visualize customers who have terminated their contracts based on international plan

```
fig, axes = plt.subplots(nrows=1, ncols=1, figsize=(15, 3))
sns.barplot(x = "churn", y = "international plan", data =
churn_international_plan, ax=axes)
axes.set_title("Churned customers by International plan");
```



- Out of the 483 customers who terminated their contracts 346 had no international plan and 137 had international plan

#Calculate the International Plan vs Churn percentage

```
International_plan_data = pd.crosstab(churn_df["international
plan"],churn_df["churn"])
International_plan_data['Percentage Churn'] =
International_plan_data.apply(lambda x : x[1]*100/(x[0]+x[1]),axis =
1)
print(International_plan_data)
```

```
churn          False  True  Percentage Churn
international plan
no              2664   346         11.495017
yes              186   137         42.414861
```

The above comparative analysis shows that:

- Out of the 3010 customers who do not have an international plan, 11.4% of customers have churned.
- Out of the 323 customers who have an international plan, 42.4% of them have terminated their accounts.

- It appears that a significant number of customers who purchased International plans are churning. This trend could possibly be attributed to connectivity issues or high call charges.


b. Analysis of churned customers based on Area Code

We shall look at the distribution of inactive customers based on their area code

```
churn_area_code = churn_df.groupby("area code")
["churn"].sum().reset_index()
churn_area_code
```

	area code	churn
0	408	122
1	415	236
2	510	125


```
fig, axes = plt.subplots(nrows=1, ncols=1, figsize=(15, 3))
sns.barplot(x = "area code", y = "churn", data = churn_area_code,
ax=axes)
axes.set_title("Churned customers by area code")
Text(0.5, 1.0, 'Churned customers by area code')
```



- The area code 415 had the most customers who terminated their contract while 408 area code had the least

c. Analysis of churn based on Voice Mail Plan

```
sns.countplot(x='voice mail plan',
               data=churn_df, hue='churn').set(title='voice mail plan
count on churn');
```




- Majority of the customers that have terminated their contract do not have voicemail plan. It could indicate that the voicemail plan might not be a highly desired or valued service among customers.

d. Analysis of churned based on Customer Service Calls

```
# Create the countplot
sns.countplot(x='customer service calls', hue='churn',
data=churn_df[churn_df['churn'] == True])

# Set the title and labels
plt.title("Number of Customer Service Calls vs. Churned")
plt.xlabel("Number of Customer Service Calls")
plt.ylabel("Count")

# Show the plot
plt.show()
```



- The above visualization shows that the majority of churned customers made 1 call to customer service. This could indicate that a significant number of customers who decided to leave the service had limited engagement with customer service, possibly suggesting that their issues or concerns were not adequately addressed.

e. Analysis of churn rates based on the different states


```
# Does different states have different churn rates?
churn_rate_state = pd.DataFrame(churn_df.groupby(["state"])
                                ['churn'].mean().sort_values(ascending = False))
print(churn_rate_state)
```

state	churn
CA	0.264706
NJ	0.264706
TX	0.250000
MD	0.242857
SC	0.233333
MI	0.219178
MS	0.215385
NV	0.212121
WA	0.212121
ME	0.209677
MT	0.205882
AR	0.200000
KS	0.185714
NY	0.180723

MN	0.178571
PA	0.177778
MA	0.169231
CT	0.162162
NC	0.161765
NH	0.160714
GA	0.148148
DE	0.147541
OK	0.147541
OR	0.141026
UT	0.138889
CO	0.136364
KY	0.135593
SD	0.133333
OH	0.128205
FL	0.126984
IN	0.126761
ID	0.123288
WY	0.116883
MO	0.111111
VT	0.109589
AL	0.100000
ND	0.096774
NM	0.096774
WV	0.094340
TN	0.094340
DC	0.092593
RI	0.092308
WI	0.089744
IL	0.086207
NE	0.081967
LA	0.078431
IA	0.068182
VA	0.064935
AZ	0.062500
AK	0.057692
HI	0.056604

visualization of the churn rates for states

```
fig, ax = plt.subplots(figsize=(20,8))
sns.barplot(x = np.linspace(0, len(churn_rate_state)-1,
len(churn_rate_state), endpoint=True),
            y = 'churn', data = churn_rate_state , ax = ax)
plt.title('Customer Churn Rate for different States', fontsize = 25)
ax.tick_params(axis = 'both', labelsize = 15)
plt.xlabel('States', fontsize = 20)
plt.ylabel('Churn Rate', fontsize = 20)
ax.set_xticklabels(churn_rate_state.index)
plt.tight_layout()
```



The visualization above shows that different states have different churn rates. California and New Jersey are the two highest churn rate states greater than 25%, while Alaska and Hawaii are the two lowest churn rate states with less than 6%.

3.4 Multivariate Analysis

a. Churn analysis - total calls vs. total charges by time period


lets visualize the performance of calls

```
features = [
    ('total day calls', 'total day charge'),
    ('total eve calls', 'total eve charge'),
    ('total night calls', 'total night charge')
]

fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(12, 6))

for i, (x, y) in enumerate(features):
    ax = axes[i] # Access the corresponding axis from the 1x3 grid
    sns.scatterplot(x=x, y=y, data=churn_df, hue='churn', ax=ax)
    ax.set_xlabel(x)
    ax.set_ylabel(y)

plt.tight_layout()
plt.show()
```



Based on the visualization above, we can draw the following observations:

- Among all the time periods, daytime calls are significantly charged higher compared to evening and nighttime calls.
- The above observations may indicate that daytime is considered a peak hour leading to higher charges
- The call charges for daytime, evening, and nighttime are higher even with fewer calls made. This may indicate that calls are also charged on duration and not necessarily the number of calls.

b. Churn analysis - total minutes vs. total charges by time period


```
# lets visualize minutes performance
```

```
features = [
    ('total day minutes', 'total day charge'),
    ('total eve minutes', 'total eve charge'),
    ('total night minutes', 'total night charge')
]
```

```
fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(12, 6))
```

```
for i, (x, y) in enumerate(features):
    ax = axes[i] # Access the corresponding axis from the 1x3 grid
    sns.scatterplot(x=x, y=y, data=churn_df, hue='churn', ax=ax)
    ax.set_xlabel(x)
    ax.set_ylabel(y)
```

```
plt.tight_layout()
plt.show()
```



- Among all the time periods, daytime minutes are significantly charged higher compared to evening and nighttime minutes.
- The above observations may indicate that daytime is considered a peak hour leading to higher charges.
- There is a linear relationship between the total minutes of daytime, evening, nighttime and the corresponding total charges. This indicates that the higher the subscription minutes the higher the charges.
- On average, customers who have terminated their accounts appear to have subscribed to more day minutes, leading to higher charges.

c. Churn analysis - total international calls and minutes vs. total international charges

```
# Lets visualize performance of international services


features = [
    ('total intl calls', 'total intl charge'),
    ('total intl minutes', 'total intl charge')
]

fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(12, 6)) # 1 row, 2 columns for the two subplots

for i, (x, y) in enumerate(features):
    ax = axes[i] # Access the corresponding axis from the 1x2 grid
    sns.scatterplot(x=x, y=y, data=churn_df, hue='churn', ax=ax)
    ax.set_xlabel(x)
    ax.set_ylabel(y)
```



```
plt.tight_layout()
plt.show()
```




- There is a linear relationship between the total international minutes and the corresponding total charges. This indicates that the higher the subscription minutes the higher the charges.
- The call charges seem to be higher even with fewer calls made. This may indicate that international calls may also be charged on duration and not necessarily the number of calls.

d. Churn analysis - total minutes and total charges

```
#total minutes during the period
churn_df['total_calls'] = (churn_df['total day calls'] +
churn_df['total eve calls'] + churn_df['total night calls'])
# total charge during the period
churn_df['total_charges'] = (churn_df['total day charge'] +
churn_df['total eve charge'] +
churn_df['total night charge'] + churn_df['total intl
charge'])
#total minutes during the period
churn_df['total_minutes'] = (churn_df['total day minutes']
+churn_df['total eve minutes'] +
churn_df['total night minutes'])

# visualization of churn performance for total minutes and charges
# Plot the lmpot
```

```
sns.lmplot(x='total_minutes', y='total_charges', data=churn_df,
hue='churn', fit_reg=False);
```




- Total minutes have a linear relationship with the total charge, indicating that as the number of minutes a customer subscribes to increases, the charge also increases.
- We can also observe that customers who have terminated their accounts tend to subscribe to higher minutes, resulting in a higher charge.

e. Churn analysis - total calls and total charges

#visualizing this performance

```
sns.lmplot(x='total_calls',
            y='total_charges',
            data=churn_df, hue='churn', fit_reg=False).set(title='Total
charges against Total Calls');
```



It is quite surprising that customers with a lower total number of calls tend to have higher charges, and a significant number of these high charges are associated with customers who have terminated their accounts.

```
#drop the comparison columns as they will not be included in our model
churn_df = churn_df.drop(columns =
['total_calls', 'total_charges', 'total_minutes'], axis=1)
```

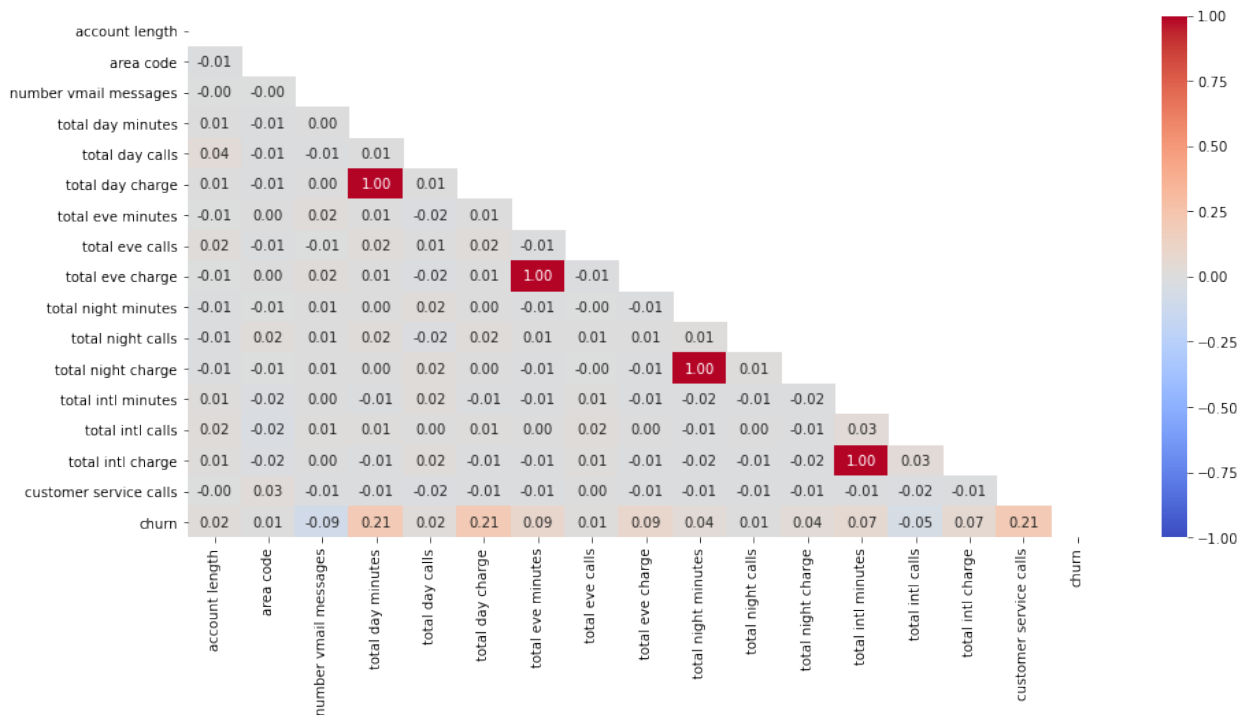
f. Visualization of Correlation Heatmap

```
# Calculate correlation matrix
corr_matrix = churn_df.corr()

# Generate a mask to hide the upper triangle
mask = np.triu(np.ones_like(corr_matrix, dtype=bool))

plt.figure(figsize=(15, 7))
sns.heatmap(corr_matrix, annot=True, vmin=-1, vmax=1, fmt=".2f",
```

```
cmap="coolwarm", mask=mask)
plt.show()
```



#checking the correlation between target variable and other features

```
churn_df.corr()['churn'].sort_values(ascending=False)
```

```
churn      1.000000
customer service calls  0.208750
total day minutes      0.205151
total day charge      0.205151
total eve minutes      0.092796
total eve charge      0.092786
total intl charge      0.068259
total intl minutes     0.068239
total night charge     0.035496
total night minutes    0.035493
total day calls        0.018459
account length         0.016541
total eve calls        0.009233
area code              0.006174
total night calls      0.006141
total intl calls       -0.052844
number vmail messages  -0.089728
Name: churn, dtype: float64
```

- ### 3.5 Preprocessing

```
churn_df_copy = churn_df.copy()
churn_df_copy
```

	number vmail messages	total day minutes	total day calls	\
0	25	265.1	110	
1	26	161.6	123	
2	0	243.4	114	
3	0	299.4	71	
4	0	166.7	113	
...	

3328	36	156.2	77
3329	0	231.1	57
3330	0	180.8	109
3331	0	213.8	105
3332	25	234.4	113

	total day charge	total eve minutes	total eve calls	total eve charge \
--	------------------	-------------------	-----------------	--------------------

0	45.07	197.4	99
16.78			
1	27.47	195.5	103
16.62			
2	41.38	121.2	110
10.30			
3	50.90	61.9	88
5.26			
4	28.34	148.3	122
12.61			
...
...			
3328	26.55	215.5	126
18.32			
3329	39.29	153.4	55
13.04			
3330	30.74	288.8	58
24.55			
3331	36.35	159.6	84
13.57			
3332	39.85	265.9	82
22.60			

	total night minutes	total night calls	total night charge \
0	244.7	91	11.01
1	254.4	103	11.45
2	162.6	104	7.32
3	196.9	89	8.86
4	186.9	121	8.41
...
3328	279.1	83	12.56
3329	191.3	123	8.61
3330	191.9	91	8.64
3331	139.2	137	6.26
3332	241.4	77	10.86

	total intl minutes	total intl calls	total intl charge \
0	10.0	3	2.70
1	13.7	3	3.70
2	12.2	5	3.29
3	6.6	7	1.78
4	10.1	3	2.73

...
3328	9.9	6	2.67
3329	9.6	4	2.59
3330	14.1	6	3.81
3331	5.0	10	1.35
3332	13.7	4	3.70

	customer service calls	churn
0	1	False
1	1	False
2	0	False
3	2	False
4	3	False
...
3328	2	False
3329	3	False
3330	2	False
3331	2	False
3332	0	False

[3333 rows x 20 columns]

Converting churn column from boolean to integer

```
churn_df_copy['churn'] = churn_df_copy['churn'].astype(int)
```

Dropping states column as it will not impact our modelling part

```
churn_df_copy = churn_df_copy.drop('state', axis=1)
```

#creating dummy variables

```
churn_df_copy= pd.get_dummies(churn_df_copy, drop_first=True)
```

```
churn_df_copy.head()
```

	account length	area code	number vmail messages	total day minutes
\				
0	128	415	25	265.1
1	107	415	26	161.6
2	137	415	0	243.4
3	84	408	0	299.4
4	75	415	0	166.7

	total day calls	total day charge	total eve minutes	total eve calls
\				
0	110	45.07	197.4	
99				
1	123	27.47	195.5	

103				
2	114	41.38	121.2	
110				
3	71	50.90	61.9	
88				
4	113	28.34	148.3	
122				

	total eve charge	total night minutes	total night calls	\
0	16.78	244.7	91	
1	16.62	254.4	103	
2	10.30	162.6	104	
3	5.26	196.9	89	
4	12.61	186.9	121	

	total night charge	total intl minutes	total intl calls	\
0	11.01	10.0	3	
1	11.45	13.7	3	
2	7.32	12.2	5	
3	8.86	6.6	7	
4	8.41	10.1	3	

	total intl charge	customer service calls	churn	international
plan_yes \				
0	2.70	1	0	
0				
1	3.70	1	0	
0				
2	3.29	0	0	
0				
3	1.78	2	0	
1				
4	2.73	3	0	
1				

	voice mail plan_yes
0	1
1	1
2	0
3	0
4	0

a. Defining the predictor and target variables

```
# define our X and y variables
X = churn_df_copy.drop (columns = ['churn'], axis=1)
y = churn_df_copy['churn']

#for consistency of results set a random seed
np.random.seed(123)
```



```

# Performing a train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.25, random_state = 42)

#scale the data
#initialize the scaler
scaler = StandardScaler()

#fit the data on the scaler
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

```

b. Fixing the class imbalance

```

# Previous original class distribution
print(y_train.value_counts())

0    2141
1     358
Name: churn, dtype: int64

# Use Smote to resample and fix the class imbalance problem
smote = SMOTE()
X_train_resampled, y_train_resampled =
smote.fit_resample(X_train_scaled, y_train)

```

We used SMOTE class in order to improve the model's performance on the minority class.

```

# Preview synthetic sample class distribution

print(pd.Series(y_train_resampled).value_counts())

1    2141
0    2141
Name: churn, dtype: int64

```

The imbalance on the target variable is now resolved.

4.0 Modeling

We will now build a model that can predict the customer churn based on the features in our dataset using the following algorithms:

- Logistic Regression
- Decision Tree
- Random Forest
- XG Boost

Model 1: Logistics Regression Classifier


```
# Instanstiate the model
logreg = LogisticRegression(random_state =42)

# fit the model
logreg.fit(X_train_resampled, y_train_resampled)

#predicting on the test
y_pred_log = logreg.predict(X_test_scaled)

def plot_confusion_matrix(y_true, y_pred, classes):
    """
    Plots a confusion matrix.
    """
    cm = confusion_matrix(y_true, y_pred)
    plt.figure()
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=classes, yticklabels=classes)
    plt.xlabel('Predicted label')
    plt.ylabel('True label')
    plt.show()

# visualizing confusion matrix
plot_confusion_matrix(y_test, y_pred_log, [0,1])
```



```
# displaying scores
print(classification_report(y_test,y_pred_log))
```

	precision	recall	f1-score	support
0	0.95	0.78	0.85	709
1	0.38	0.78	0.51	125
accuracy			0.78	834
macro avg	0.67	0.78	0.68	834
weighted avg	0.87	0.78	0.80	834

Logistics Regression observations

Recall measures the ability of the model to correctly identify customers who are likely to churn (positive instances) out of all the customers who actually churned.

- For class 0, which represents customers who did not churn, the recall is 0.78. This means that the model correctly identified 78% of the customers who did not churn out of the total number of customers who actually did not churn.
- Similarly, for class 1, which represents customers who churned, the recall is 0.78, indicating that the model correctly identified 78% of the customers who churned out of the total number of customers who actually churned.

Accuracy: 0.78 means that 78% of the total number of customers was correctly classified.


Model 2: Decision Tree Classifier

```
# Instantiate a DT classifier
clf = DecisionTreeClassifier(random_state=42)

# fit DT classifier
clf.fit(X_train_resampled, y_train_resampled)

# Make predictions for test data
y_pred_clf = clf.predict(X_test_scaled)

# plotting a confusion matrix
plot_confusion_matrix(y_test, y_pred_clf, [0,1])
```



```
print(classification_report(y_test,y_pred_clf))
```

	precision	recall	f1-score	support
0	0.96	0.93	0.94	709
1	0.64	0.76	0.70	125
accuracy			0.90	834
macro avg	0.80	0.84	0.82	834
weighted avg	0.91	0.90	0.90	834

Decision Tree Classifier Observations

Recall

- For class 0, which represents customers who did not churn, the recall is 0.93. This means that the model correctly identified 93% of the customers who did not churn out of the total number of customers who actually did not churn.
- Similarly, for class 1, which represents customers who churned, the recall is 0.76, indicating that the model correctly identified 76% of the customers who churned out of the total number of customers who actually churned.

Accuracy: 0.90 means that 90% of the total number of customers was correctly classified. The model performs better than logistics regression model.


Model 3: Random Forest Classifier

```
# Instantiate a DT classifier
rfc = RandomForestClassifier(random_state=42)
```

```
# fit RFCclassifier
rfc.fit(X_train_resampled, y_train_resampled)

# Make predictions for test data
y_pred_rfc = rfc.predict(X_test_scaled)

plot_confusion_matrix(y_test, y_pred_rfc, [0,1])
```



```
print(classification_report(y_test,y_pred_rfc))
```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	709
1	0.84	0.82	0.83	125
accuracy			0.95	834
macro avg	0.91	0.89	0.90	834
weighted avg	0.95	0.95	0.95	834

Random Forest Classifier Observations

Recall

- For class 0, which represents customers who did not churn, the recall is 0.97. This means that the model correctly identified 97% of the customers who did not churn out of the total number of customers who actually did not churn.

- Similarly, for class 1, which represents customers who churned, the recall is 0.82, indicating that the model correctly identified 82% of the customers who churned out of the total number of customers who actually churned.

Accuracy: 0.95 means that 95% of the total number of customers was correctly classified. The model performs better than Decision Tree Classifier model.


Model 4:XGBoost

```
# Instantiate the model
x_gb = XGBClassifier(random_state=42)

# fit XGB classifier
x_gb.fit(X_train_resampled, y_train_resampled)

# Make predictions for test data
y_pred_xgb = x_gb.predict(X_test_scaled)

plot_confusion_matrix(y_test, y_pred_xgb, [0,1])
```



```
print(classification_report(y_test,y_pred_xgb))
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	709
1	0.92	0.81	0.86	125
accuracy			0.96	834
macro avg	0.94	0.90	0.92	834

weighted avg	0.96	0.96	0.96	834
--------------	------	------	------	-----

XGBoost Classifier Observations

Recall

- For class 0, which represents customers who did not churn, the recall is 0.99. This means that the model correctly identified 99% of the customers who did not churn out of the total number of customers who actually did not churn.
- Similarly, for class 1, which represents customers who churned, the recall is 0.81, indicating that the model correctly identified 81% of the customers who churned out of the total number of customers who actually churned.

Accuracy: 0.96 means that 96% of the total number of customers was correctly classified. The model seems to perform as well as the Random Forest Classifier model.

5.0 Model Evaluation

5.1 Model comparison

```
classifiers = [LogisticRegression(),
                RandomForestClassifier(),
                DecisionTreeClassifier(),
                XGBClassifier()]

# Define a result table as a DataFrame
result_table = pd.DataFrame(columns=['classifiers', 'accuracy',
                                     'recall'])

# Train the models and record the results
for cls in classifiers:
    model = cls.fit(X_train_resampled, y_train_resampled)
    y_pred = model.predict(X_test_scaled)
    accuracy = accuracy_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred) # Calculate precision
    score
    result_table = result_table.append({'classifiers':
cls.__class__.__name__,
                                     'accuracy': accuracy,
                                     'recall': recall}, ignore_index=True)

# Set name of the classifiers as index labels
result_table.set_index('classifiers', inplace=True)

result_table
```

	accuracy	recall
classifiers		
LogisticRegression	0.775779	0.776
RandomForestClassifier	0.954436	0.832
DecisionTreeClassifier	0.894484	0.744
XGBClassifier	0.960432	0.808

- All the models are able to predict well, however, Random Forest Classifier and XGBoost Classifier have the highest accuracy and recall scores.
- We shall proceed and tune Random Forest Classifier and XGBoost classifier hyperparameters and compare the results.


ROC

```
# Get the ROC curves for all classifiers
classifiers = ["LogisticRegression", "DecisionTreeClassifier",
               "RandomForestClassifier", "XGBClassifier"]
roc_curves = []

for classifier_name in classifiers:
    if classifier_name == "LogisticRegression":
        classifier = LogisticRegression()
    elif classifier_name == "DecisionTreeClassifier":
        classifier = DecisionTreeClassifier()
    elif classifier_name == "RandomForestClassifier":
        classifier = RandomForestClassifier()
    elif classifier_name == "XGBClassifier":
        classifier = XGBClassifier()

    classifier.fit(X_train_resampled, y_train_resampled)
    y_pred = classifier.predict(X_test_scaled)
    fpr, tpr, _ = roc_curve(y_test, y_pred)
    roc_auc = auc(fpr, tpr)
    roc_curves.append((fpr, tpr, roc_auc, classifier_name))

# Plot the ROC curves and print AUC values
plt.figure()
for fpr, tpr, roc_auc, classifier_name in roc_curves:
    plt.plot(fpr, tpr, label=f'{classifier_name} (AUC = {roc_auc:.2f})')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.show()
```

- XGB Classifier and RandomForestClassifier are producing better results in model 4 and model 3 respectively.
- The AUC value for model 3: RandomForest is 0.90 and Model 4: XGBoost is 0.90
- Lets perform hyperparameter tuning to improve them.

5.2 Hyperparameter tuning for our best models

1. Tuned RandomForestClassifier`

```
# Create a parameter grid with reduced values
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [5, 10],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2],
}

# Create a grid search object
rfc = RandomForestClassifier()
grid_search = GridSearchCV(rfc, param_grid, cv=3, scoring='accuracy',
n_jobs=-1)

# Fit the grid search object
grid_search.fit(X_train_resampled, y_train_resampled)

# Print the best parameters
print(grid_search.best_params_)
```

```
{'max_depth': 10, 'min_samples_leaf': 1, 'min_samples_split': 2,
'n_estimators': 100}
```

```
# Instantiate a RandomForest classifier
```

```
rfc_tune = RandomForestClassifier(max_depth=10,
                                min_samples_leaf=1,
                                min_samples_split=2,
                                n_estimators=200,
                                random_state=42)
```


```
# fit RFCclassifier
```

```
rfc_tune.fit(X_train_resampled, y_train_resampled)
```

```
# Make predictions for test data
```

```
y_pred_rfc_tune = rfc_tune.predict(X_test_scaled)
```

```
plot_confusion_matrix(y_test, y_pred_rfc_tune, [0,1])
```



```
print(classification_report(y_test,y_pred_rfc_tune))
```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	709
1	0.83	0.82	0.82	125
accuracy			0.95	834
macro avg	0.90	0.89	0.90	834
weighted avg	0.95	0.95	0.95	834


ROC curve

```
# Make predictions
y_prob = x_gb.predict(X_test_scaled)
y_pred = (y_prob > 0.5).astype(int)

# Calculate evaluation metrics
roc_value = roc_auc_score(y_test, y_prob)
print("RNN roc_value: {0}" .format(roc_value))
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
threshold = thresholds[np.argmax(tpr-fpr)]

roc_auc = auc(fpr, tpr)
print("ROC for the test dataset", '{:.1%}' .format(roc_auc))
plt.plot(fpr, tpr, label="Test, auc="+str(roc_auc))
plt.legend(loc=4)
plt.show()

RNN roc_value: 0.8976530324400565
ROC for the test dataset 89.8%
```



Checking for Overfitting

```
# Make predictions for test data
y_train_pred_rfc = rfc_tune.predict(X_train_resampled)
y_test_pred_rfc = rfc_tune.predict(X_test_scaled)
```

```
# Calculate accuracy on the training and test data
train_accuracy = accuracy_score(y_train_resampled, y_train_pred_rfc)
test_accuracy = accuracy_score(y_test, y_test_pred_rfc)

print("Train Accuracy:", train_accuracy)
print("Test Accuracy:", test_accuracy)

Train Accuracy: 0.9673049976646427
Test Accuracy: 0.947242206235012
```

Tuned Random Forest Classifier Observations

Recall

- For class 0, which represents customers who did not churn, the recall is 0.97. This means that the model correctly identified 97% of the customers who did not churn out of the total number of customers who actually did not churn.
- Similarly, for class 1, which represents customers who churned, the recall is 0.82, indicating that the model correctly identified 82% of the customers who churned out of the total number of customers who actually churned.

Accuracy: 0.95 means that 95% of the total number of customers was correctly classified. The model performs better than Decision Tree Classifier model.

Important Features for RandomForest Model

```
# Assuming 'churn' is the target column, and you want to remove it
from churn_df_copy
# You can create a new DataFrame without the 'churn' column
churn_df_copy_without_churn = churn_df_copy.drop('churn', axis=1)


# Get the feature importances from the XGBoost model
importances = rfc_tune.feature_importances_

# Get the indices to sort the features in descending order of
importance
indices = np.argsort(importances)[::-1]

# Get the feature names and importances for the top 10 features
top_n = 10
top_feature_names =
churn_df_copy_without_churn.columns[indices[:top_n]]
top_importances = importances[indices[:top_n]]

# Plot the top 10 feature importances as a horizontal bar plot
plt.figure(figsize=(8, 6))
plt.barh(range(top_n), top_importances, align='center')
plt.yticks(range(top_n), top_feature_names)
plt.xlabel('Relative Importance')
plt.ylabel('Features')
```

```
plt.title('Top 10 Feature Importances in RandomForestClassifier Model')
plt.show()
```



According to the Random Forest Model, customer service calls, total day charge and international plan yes are the top 3 most important features contributing to customer churn.

2. Tuned XGBoost Classifier

```
parameters = {
    'max_depth': range(3,10,2),
    'min_child_weight': range(1,6,2),
    'gamma': [i/10.0 for i in range(0,5)],
    'learning_rate' : [i/10.0 for i in range(0,5)],
    'n_estimators': range(10,150,10)
}
random_search=RandomizedSearchCV(estimator =
XGBClassifier(base_score=0.5, booster='gbtree',
colsample_bylevel=1,colsample_bynode=1, colsample_bytree=1, gamma=0,
learning_rate=0.1, max_delta_step=0, max_depth=3,min_child_weight=1,
n_estimators=100, n_jobs=-1,
nthread=None, objective='binary:logistic',
```

```

random_state=42,reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
seed=None,

silent=None, subsample=1, verbosity=1),
param_distributions=parameters,scoring='roc_auc',n_jobs=4,cv=5)
random_search.fit(X_train_resampled, y_train_resampled)
random_search.best_params_

{'n_estimators': 120,
 'min_child_weight': 5,
 'max_depth': 9,
 'learning_rate': 0.3,
 'gamma': 0.0}


# Instanstiate the model
x_gb_tune = XGBClassifier(learning_rate=0.3, max_depth=9,
                          n_estimators=120, min_child_weight = 5,
                          gamma = 0.0, random_state=42)

# fit XGB classifier
x_gb_tune.fit(X_train_resampled, y_train_resampled)

# Make predictions for test data
y_pred_xgb_tune = x_gb.predict(X_test_scaled)

# Plotting confusion matrix
plot_confusion_matrix(y_test, y_pred_xgb_tune, [0,1])

```



```
# display scores
```

```
print(classification_report(y_test,y_pred_xgb_tune))
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	709
1	0.92	0.81	0.86	125
accuracy			0.96	834
macro avg	0.94	0.90	0.92	834
weighted avg	0.96	0.96	0.96	834

ROC

```
# Make predictions
```

```
y_prob = x_gb.predict(X_test_scaled)
```

```
y_pred = (y_prob > 0.5).astype(int)
```

```
# Calculate evaluation metrics
```

```
roc_value = roc_auc_score(y_test, y_prob)
```

```
print("RNN roc_value: {0}" .format(roc_value))
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
```

```
threshold = thresholds[np.argmax(tpr-fpr)]
```

```
roc_auc = auc(fpr, tpr)
```

```
print("ROC for the test dataset", '{:.1%}'.format(roc_auc))
```


```
plt.plot(fpr, tpr, label="Test, auc="+str(roc_auc))
```

```
plt.legend(loc=4)
```

```
plt.show()
```

```
RNN roc_value: 0.8976530324400565
```

```
ROC for the test dataset 89.8%
```



Checking for Overfitting

```
# Make predictions for test data
y_train_pred_xgb = x_gb.predict(X_train_resampled)
y_test_pred_xgb = x_gb.predict(X_test_scaled)

# Calculate accuracy on the training and test data
train_accuracy = accuracy_score(y_train_resampled, y_train_pred_xgb)
test_accuracy = accuracy_score(y_test, y_test_pred_xgb)

print("Train Accuracy:", train_accuracy)
print("Test Accuracy:", test_accuracy)
```

```
Train Accuracy: 1.0
Test Accuracy: 0.960431654676259
```

Tuned XGBoost Classifier Observations

Recall

- For class 0, which represents customers who did not churn, the recall is 0.99. This means that the model correctly identified 99% of the customers who did not churn out of the total number of customers who actually did not churn.
- Similarly, for class 1, which represents customers who churned, the recall is 0.81, indicating that the model correctly identified 81% of the customers who churned out of the total number of customers who actually churned.

Accuracy: 0.96 means that 96% of the total number of customers was correctly classified. The model performs better than Decision Tree Classifier model.

Important Features for tuned XGBoost Model


```
# Assuming 'churn' is the target column, and you want to remove it
from churn_df_copy
# You can create a new DataFrame without the 'churn' column
churn_df_copy_without_churn = churn_df_copy.drop('churn', axis=1)

# Get the feature importances from the XGBoost model
importances = x_gb_tune.feature_importances_

# Get the indices to sort the features in descending order of
importance
indices = np.argsort(importances)[::-1]

# Get the feature names and importances for the top 10 features
top_n = 10
top_feature_names =
churn_df_copy_without_churn.columns[indices[:top_n]]
top_importances = importances[indices][:top_n]

# Plot the top 10 feature importances as a horizontal bar plot
plt.figure(figsize=(8, 6))
plt.barh(range(top_n), top_importances, align='center')
plt.yticks(range(top_n), top_feature_names)
plt.xlabel('Relative Importance')
plt.ylabel('Features')
plt.title('Top 10 Feature Importances in XGBoost Model')
plt.show()
```



According to the XGBoost Model, international plan yes, customer service calls and voice mail plan_yes are the top 3 most important features contributing to customer churn.

Conclusion

- **RandomForestClassifier**
 - Tuned RandomForest model with an AUC of 0.89 suggesting that the model has a strong ability to distinguish between positive (churned) and negative (not churned) instances.
 - This indicates that the model has a good balance between sensitivity (recall) and specificity, capturing a high proportion of both churned and non-churned customers accurately.
 - The recall values varied slightly, with the Tuned Random Forest Classifier performing slightly better in identifying customers who churned at 82%.
 - Tuned Random Forest Classifier had an accuracy of 95% of the total number of customers that were correctly classified
- **XGBoost Classifier**
 - Tuned XGBoost Classifier had an AUC of 0.89, it had a recall for class 1 at 81% and an accuracy of 96%
- **Picking the best model**
 - After carefully analyzing the performance metrics of both models, the Tuned Random Forest Classifier emerges as the better choice for our specific objective of correctly identifying churn customers. With a recall of 82%, the

model accurately identifies 82% of the customers who churned out of the total number of customers who actually churned.

- While the Tuned XGBoost Classifier exhibits a slightly higher accuracy (96%) and recall for non-churn customers (99%), our primary focus lies in correctly identifying churn customers to effectively target retention strategies. The Tuned Random Forest Classifier's recall of 82% for churn customers is commendable and aligns better with our priority.
- Therefore, we confidently select the Tuned Random Forest Classifier as our best model for predicting customer churn and enabling us to take proactive measures to retain valuable customers, thereby enhancing overall business performance.
- Based on the analysis using our best model (Random Forest Classifier), we can confidently conclude that the three most important factors influencing churn are the number of customer service calls made, the total day charge incurred, and the presence of an international plan

Summary Findings

- Majority of customers who terminated their contracts did not have a voicemail plan.
- California and New Jersey have the highest churn rates, both exceeding 25%
- Customers who terminated their accounts appeared to have subscribed to more day minutes, resulting in higher charges.
- Charges for total daytime calls and minutes were significantly higher compared to evening and nighttime calls and minutes.
- There is a lack of proportionality between the total number of international calls made and the corresponding charges, meaning that charges are higher even with fewer total calls.
- The customers with international plan have the higher churn rate compared to those with no plan.

Recommendation

- Ensure fairness in charging, establish a proportional charge for daytime, evening, nighttime and international calls.
- Enhance the voice mail plan service to be more appealing to customers.
- Bring down cost of daytime calls and minutes charges
- Focus more on customer service